

“When everything can look real, trust becomes the product.”

Prod-A-Thon

Designing **Real-Time Trust Signals** in a **Near-Perfect AI Visual Feed**

By : Team WTH

Aarya Shisode, Akshada
Bhamare, Deepanjali

Institute Name : IIT Guwahati

Understanding the problem

Introduction & Current Landscape

Instagram’s Mission 

To help people **connect** and **share** safely while making sure they can **trust** and **understand** the content they see, especially as **AI-made** media gets more realistic.

As **AI tools for image and video creation rapidly improve**, it has become increasingly difficult for users to distinguish **real moments** from **altered or fully synthetic content**. Current systems rely on **reactive labels and post-hoc reporting**, which appear only **after users have already consumed or engaged with the content**.

Key Assumptions

AI content will grow exponentially

Users want trust without feed disruption

Instagram can combine metadata, disclosure, and signals

Creators will disclose if it’s easy and fair

False positives must be near zero

What is the problem at hand ?

Authenticity Blind Spot: Users face an authenticity blind spot-on Instagram, struggling to distinguish between real captures, edited images, and fully AI-generated content as visual quality converges.

Reactive & Confusing Labels

- Current “AI info” labels appear after engagement
- Binary labeling fails to reflect degree of manipulation

Trust Erosion

- Users increasingly question the credibility of visual content
- Over-labeling leads to skepticism even toward real photos

Gaps Identified 

1


Different authenticity signals (metadata, disclosures, detection) are not unified, leaving users without a single source of truth.

2

Users lack a quick, intentional way to verify content before liking, sharing, or believing it.

3

Labels don’t explain what was changed, leading to confusion, misinterpretation, and reduced trust.



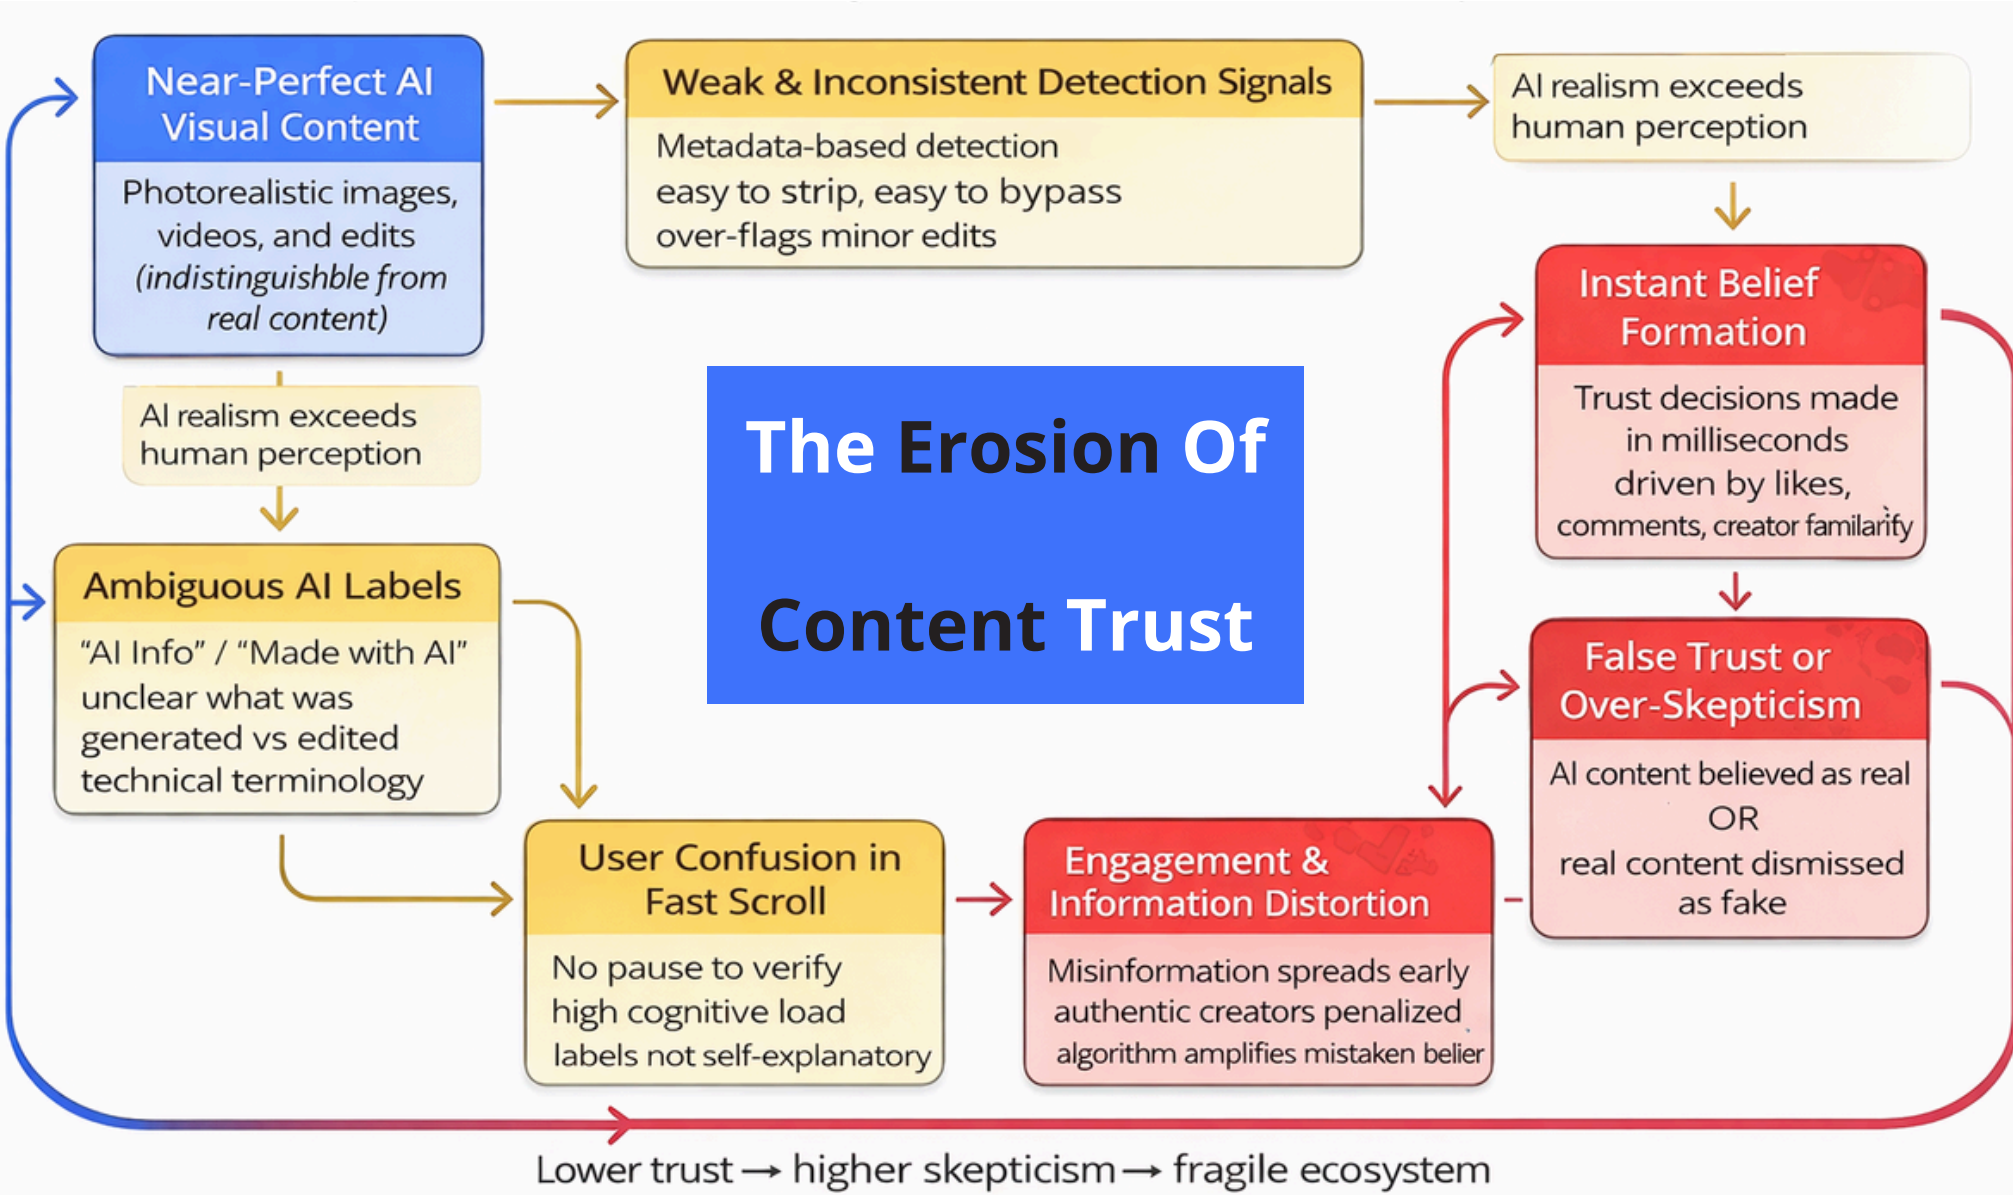
Need for AI Detection and Benchmarking

 **Users can’t detect AI:** AI images are identified correctly only 62% of the time and AI videos 23%, while detection tools perform at 50–65%. Trust forms instantly during scroll, driven by social proof.

 **Trust breaks at scale:** 60% of users doubt online authenticity and 62% avoid AI-labeled content. One misleading post erodes trust across the entire feed.

 **Current systems fail:** Metadata is easy to bypass, AI video/audio lacks detection, and false positives cause 15–80% engagement loss, despite needing <1% false positives. Labels arrive after belief forms.

How AI-Generated Content Creates a Trust Breakdown Loop on Instagram



Instagram User Base: Segmentation

Dimension	Key Finding	Implication
Human Detection (Images)	~62% accuracy	Users misjudge AI images frequently
Human Detection (Videos)	~23% accuracy	AI videos are almost indistinguishable
AI Detectors (Real Feeds)	50–65% accuracy	Automated detection is unreliable
Content Alterations	Filters & compression reduce signals	Labels break after reposts
Trust Impact	Loss spreads feed-wide	Implication

The Trust-at-Glance Effect Visual realism bias: AI visuals match or exceed real photography, triggering automatic belief. Social proof shortcuts: Likes, comments, and creator familiarity replace authenticity judgment. Emotional hooks: Beauty, aspiration, and novelty reduce critical evaluation. Feed velocity pressure: Fast scrolling forces snap decisions over conscious verification. Key Insight: Trust is formed in seconds, before users even consider authenticity.	The Authenticity Blind Zone Invisible manipulation: AI edits and filters leave no obvious visual traces. Binary label confusion: “AI info” feels absolute even for minor or partial edits. Metadata dependency: Detection relies on tags, not actual visual manipulation. Metadata dependency: Detection relies on tags, not actual visual manipulation. Key Insight: Users are forced to guess authenticity with incomplete and confusing signals.	The Slow Trust Erosion Loop Normalization of unreal: AI-perfect visuals reset expectations of reality. Self-doubt in judgment: Users lose confidence in their ability to tell what’s real. Creator skepticism spillover: Even genuine creators face disbelief. Platform trust decay: One bad experience reduces trust across the entire feed. Key Insight: Authenticity confusion compounds silently, eroding trust over time
--	---	---

Trust & Authenticity Ecosystem

- Humans detect AI images with only 62% accuracy
- AI videos detected correctly just 23% of the time
- Belief forms before verification
- Social proof drives trust ($p < 0.001$)
- AI detectors drop to 50–65% accuracy in real feeds
- Accuracy degrades after filters, compression, edits
- False positives disproportionately affect non-native creators
- <1% false positive rate is non-negotiable
- Trust loss often begins after just one misleading exposure, with damage forming instantly during the scroll—well before users report, verify, or fact-check the content. By the time moderation intervenes, perception is already shaped, causing trust erosion that spreads across the entire feed rather than remaining confined to a single post.

User Pain points

Unmeet Needs



“I’m Aanya (21), a college student, and I spend most of my time on Instagram Reels. I follow beauty and lifestyle influencers and want to scroll without constantly questioning whether what I’m seeing is real or AI-generated.”



Instant authenticity cues ($\leq 1s$) that reassure what’s real or AI without pausing the feed, while reducing self-comparison without shaming creators.

Why Authenticity Fails



- There are no instant, in-feed authenticity cues to judge real versus AI content while scrolling.
- AI, filters, and edits appear visually identical, making authenticity impossible to assess at a glance.
- Trust is formed emotionally within seconds, long before any verification mechanism becomes visible.
- Verification requires extra effort or stopping the feed, so use skip it entirely.

2.The Intent-Driven Viewer



“I’m Rohan (29), a working professional, and I use Instagram to discover travel spots, cafés, and products. Before spending money or planning a trip, I want to know whether what I’m seeing reflects reality or is heavily edited or AI-generated.”



Quick credibility checks before acting, with a clear enhanced vs synthetic distinction, delivered through non-interruptive trust signals.



- There is no quick way to verify authenticity before making purchase, travel, or lifestyle decisions.
- Enhanced visuals and fully AI-generated content are mixed together without clear differentiation.
- Captions, disclosures, or reporting tools appear too late to influence decision-making.
- Distracted browsing during short sessions leaves no time for deeper authenticity checks.

3.The Credibility-Seeking Creator



“I’m Neeraj (32), a photographer and content creator. I capture real moments, but my work is often questioned as ‘AI-generated’ because it looks polished. I want audiences and brands to trust my authenticity without hurting my reach.”




Quick credibility checks before acting, with a clear enhanced vs synthetic distinction, delivered through non-interruptive trust signals.



- There is no in-feed mechanism that allows creators to prove the originality of their content.
- High-quality authentic photography is increasingly mistaken for AI-generated or synthetic visuals.
- Creators lack voluntary, creator-controlled transparency tools to signal authenticity safely.
- Fear of mislabeling creates anxiety around reach, partnerships ,and long-term credibility


Solution-



Instagram adds a **lightweight, scroll-speed trust system** that helps users judge authenticity **without interrupting the feed**. Posts show **subtle confidence indicators** based on metadata, creator disclosure, and high-confidence AI detection. Users can tap **“Verify Authenticity”** for a **layered explanation**: a quick confidence snapshot, a plain-language summary of edits or AI use, and optional technical proof. Creators are protected through **voluntary AI disclosures and authenticity profiles**, reducing false positives.

In essence: Instagram empowers users with **confidence, transparency, and control**—instantly and unobtrusively.

Instagram Authenticity Confidence Stack



The Authenticity Confidence Stack is designed to work at scroll speed, adding trust **without disrupting the feed experience**. Instead of binary “real vs fake” labels, Instagram computes an **Authenticity Confidence Score** using multiple high-precision signals—visual artifacts, social-media-robust detection, metadata, and creator disclosures. Only **high-certainty cases** surface indicators, minimizing false positives.

At the **feed level**, users see **subtle, non-judgmental signals** (e.g., AI-assisted, edited) that communicate how confidently Instagram understands a post’s origin—without accusations or warnings.

For users seeking clarity, “Verify Authenticity” opens a layered explanation:

- **A quick confidence snapshot**
- **Plain-language explanation of edits or AI use**
- **Optional technical details for advanced users**

Creators are protected through voluntary AI disclosure tools and authenticity profiles, ensuring real creators are not penalized and edge cases remain contestable.

A horizontal slider bar with five segments: Problem, Stakeholder, Solution, Implementation, and Metrics & Pitfall. The 'Solution' segment is highlighted in blue, indicating it is the current focus.

A screenshot of the Instagram profile page for 'photoreanher'. The profile picture is a circular icon with a camera lens. The menu is open, showing the following options: 'Report', 'About This Account', 'Copy Link', 'Turn Off Commenting' (with a toggle switch turned on), 'Share, or Deampost...', and 'Verify Authenticity' (highlighted with a blue checkmark icon). The 'OK' button is at the bottom of the menu.

The screenshot shows the 'Verify Authenticity' interface in the Adobe Lightroom mobile app. At the top, there's a navigation bar with a back arrow and a three-dot menu. Below this is the profile header for 'Verify Authenticity' with a camera icon and the handle '@photographer'. The main section is titled 'Edited Photo'. A summary bar indicates '91% High confidence' with a green 'Relet' button and a chevron. Below this are three rows of detected edits: 'Captured on phone', 'Lighting adjusted in Lightroom', and 'Minor retouching applied'. At the bottom, there is a link for 'Optional technical details' and a large 'OK' button.

< New Post Share

@creator Authenticity Signals >

☒ AI-Assisted

Prexy.ey inloyore

Phone Phone
15 Pro, Lightroom

Next

Instagram Authenticity Confidence Stack

The diagram illustrates the process of verifying a post's authenticity on Instagram. It starts with a 'Feed scroll' and 'Passive authenticity signals' leading into a 'Passive signals' box. This box contains three items: 'Source Verified / AI-Assisted' (with a checkmark icon), 'Lighting adjusted' (with a sun icon), and 'Minor retouching applied' (with a star icon). From the 'Passive signals' box, the flow goes to an 'Optional technical details' box. Both the 'Passive signals' and 'Optional technical details' boxes lead to a 'Request Verified' state. The 'Request Verified' state includes three items: 'Disclosed' (with a checkmark icon), 'Device 15 Pro, Lightroom' (with a camera icon), and 'High' (with a star icon and a 'P' icon). A dashed arrow indicates a feedback loop from the 'Request Verified' state back to the 'Passive signals' box.

```

graph LR
    FS[Feed scroll] --> PS[Passive signals]
    PAS[Passive authenticity signals] --> PS
    PS --> OTD[Optional technical details]
    OTD --> RV[Request Verified]
    PS --> RV
    RV -.-> PS
  
```

Passive signals

- ✓ Source Verified / AI-Assisted
- ☀ Lighting adjusted
- ★ Minor retouching applied


Optional technical details

Request Verified

- ✓ Disclosed >
- 📷 Device 15 Pro, Lightroom
- 🌟 High ★★★★★ P

- Creator disclosures + system signals → authenticity confidence
- User verification & feedback → model improvement
- High-confidence outcomes only → in-feed signals

Platform Performance & Risk Management



Metrics

North Star Metric

Informed Trust Rate (ITR)

% of content interactions (views, shares, saves) where users engage after seeing an authenticity signal or verification layer.

Why this matters:

- Measures **confidence**, not blind consumption
- Aligns with Instagram's **trust + engagement** goals
- Avoids **over-optimizing** for detection accuracy alone

Success Metrics (AARRR Framework)

Activation – Authenticity Interaction Rate

% of users who interact with authenticity signals at least once

A

Activation – Creator Disclosure Completion Rate

- % of creators completing AI disclosure during upload
- Indicates adoption of transparency tools

A

Retention – Trust-Aware MAU Ratio

- % of MAUs who repeatedly engage with verified or disclosed content
- Shows sustained trust, not novelty usage

R

Referral – Share-with-Context Rate

- % of shares where authenticity context is viewed before sharing
- Reduces misinformation spread velocity

R

Revenue – Brand Safety Confidence Lift

- Increase in brand-safe impressions and creator-brand partnerships
- Trust → monetization flywheel

R

Pitfalls & Mitigations

	1	2	3	4	5
Main Risk	Over-Labeling & Trust Fatigue	Model Bias & Detection Errors	User Misinterpretation of Labels	Creator Backlash & Reach Anxiety	Low Feature Adoption
Probability	○○○○○	○○○○○	○○○○○	○○○○○	○○○○○
Impact	○○○○○	○○○○○	○○○○○	○○○○○	○○○○○
Mitigation	- Passive signals only - Labels shown only at >95% confidence - No “fake” language	- Precision-first thresholds - Multi-signal verification (metadata + models + disclosure)	- Plain-language explanations - Progressive disclosure - One-time education tooltips	- No reach penalties for disclosure - Authenticity badge rewards transparency - Preview of labels before posting	- Zero disruption to feed - Optional deep dive - Trust signals integrated into familiar UI

Our Strategic Edge




Real-Time Authenticity Scoring

- Surface confidence during feed ranking
- Reduce misinformation spread before virality
- Shift from post-hoc to pre-belief trust



Federated Learning for Detection

- Improve AI detection without sharing user data
- Improve robustness across regions



Cross-Platform Provenance

- Shared authenticity credentials across Meta platforms
- Consistent trust signals on Instagram, Threads, Facebook

Prod-A-Thon

Thank You

By : Team WTH

**Aarya Shisode,Akshada
Bhamare,Deepanjali**

Institute Name : IIT Guwahati